

# gpu-fryer summary

- Generated: 2026-08-06 17:19:05
- Nodes: node5500 (8 x B200), node5501 (8 x B200), node5502 (8 x B200)
- Note: node5500-5502 are no longer in Slurm. The tables below are their historical baseline; the current hardware is the 7 new nodes covered in "New nodes" below.
- Precisions: FP32, BF16, FP8
- Reference (MIT aicr-benchmarks, gpu-fryer/summary.md, b0025, **B200**): per-GPU mean TFLOP/s — FP32 772, BF16 1500, FP8 4115

## Per-node mean converged throughput (TFLOP/s)

| Node              | GPU  | FP32 | BF16 | FP8  | Health |
|-------------------|------|------|------|------|--------|
| node5500          | B200 | 744  | 1457 | 3990 | ok     |
| node5501          | B200 | 748  | 1457 | 4001 | ok     |
| node5502          | B200 | 760  | 1437 | 4062 | ok     |
| reference (b0025) | B200 | 772  | 1500 | 4115 | —      |

## % of B200 reference (mean, B200 nodes only)

| Node     | FP32 | BF16 | FP8 |
|----------|------|------|-----|
| node5500 | 96%  | 97%  | 97% |
| node5501 | 97%  | 97%  | 97% |
| node5502 | 98%  | 96%  | 99% |

## New nodes (node5600/5601/5602/5702/5800/5801/5802)

Run 2026-08-24 on the 7 new B200 nodes, same configuration as above (300 s per precision, all 8 GPUs, 7 single-node jobs).

**They match the nodes above — no separate tables needed.** Per-node means land at FP32 740-752, BF16 1451-1465, FP8 3949-4032 TFLOP/s. Against the node5500-5502 means the 7-node means differ by **-0.4% (FP32)**, **+0.5% (BF16)**, **-0.3% (FP8)**, and no individual node is more than **1.7%** away on any precision (worst: node5600-c1, FP8). Every node lands at 96-98% of the b0025 B200 reference, the same band the old nodes occupied.

**All 7 are healthy.** gpu-fryer reports "All GPUs seem healthy" on every node;  
no HW, thermal-SW or thermal-HW throttling flag was raised on any of the 56 GPUs.  
Per-GPU uniformity is also unchanged: at 8 GPUs the derived speed-up is  
7.81-7.88x, against 7.87-8.03x on the old nodes.

## Speed-up vs number of GPUs

![[Speedup vs number of GPUs](gpu-fryer-speedup.svg)]

Shown for **node5500**, one curve per precision.

| #GPUs | BF16 | FP32 | FP8  | ideal |
|-------|------|------|------|-------|
| 1     | 1.00 | 1.00 | 1.00 | 1.00  |
| 2     | 1.99 | 2.02 | 2.00 | 2.00  |
| 3     | 2.97 | 3.02 | 2.97 | 3.00  |
| 4     | 3.94 | 4.01 | 3.95 | 4.00  |
| 5     | 4.93 | 5.03 | 4.93 | 5.00  |
| 6     | 5.92 | 6.03 | 5.93 | 6.00  |
| 7     | 6.90 | 7.03 | 6.90 | 7.00  |
| 8     | 7.87 | 8.03 | 7.88 | 8.00  |

The other nodes (node5501, node5502) are close and are omitted from the plot to keep it readable: at 8 GPUs node5500 reaches BF16 7.87x, FP32 8.03x, FP8 7.88x, and no other node differs by more than **0.16x** on any precision.

> **How to read this.** gpu-fryer stresses all 8 GPUs *\*concurrently\** and reports one converged figure per GPU — it does not run separate 1, 2, ... 8-GPU jobs. The curve above is therefore **derived** from that single run:  
 $\text{speed-up}(N) = (\text{sum of GPUs } 0..N-1) / \text{GPU } 0$ . It is linear by construction and is **not** a measured scaling study; what it shows is per-GPU *\*uniformity\** — a curve that tracks the dashed ideal line means every GPU sustains the same throughput, while a curve bending below it marks a slow or throttling GPU. For real scaling behaviour see the Megatron-LM weak-scaling results in `output-megatron/summary.md`.

## Per-GPU converged throughput (TFLOP/s)

### node5500 (8 x B200)

| GPU | FP32  | BF16   | FP8    |
|-----|-------|--------|--------|
| 0   | 740.9 | 1480.4 | 4049.2 |
| 1   | 755.1 | 1470.5 | 4064.6 |
| 2   | 738.5 | 1440.6 | 3921.7 |

|             |              |               |               |
|-------------|--------------|---------------|---------------|
| 3           | 736.1        | 1439.6        | 3940.8        |
| 4           | 752.7        | 1470.3        | 4002.6        |
| 5           | 747.9        | 1458.6        | 4035.6        |
| 6           | 740.8        | 1450.2        | 3920.9        |
| 7           | 740.9        | 1444.2        | 3988.4        |
| <b>min</b>  | <b>736.1</b> | <b>1439.6</b> | <b>3920.9</b> |
| <b>mean</b> | <b>744.1</b> | <b>1456.8</b> | <b>3990.5</b> |
| <b>max</b>  | <b>755.1</b> | <b>1480.4</b> | <b>4064.6</b> |

### node5501 (8 x B200)

| GPU         | FP32         | BF16          | FP8           |
|-------------|--------------|---------------|---------------|
| 0           | 759.2        | 1476.6        | 4059.5        |
| 1           | 747.3        | 1452.8        | 4012.1        |
| 2           | 742.4        | 1446.1        | 3993.2        |
| 3           | 754.4        | 1469.3        | 4026.3        |
| 4           | 735.4        | 1429.2        | 3974.8        |
| 5           | 749.6        | 1460.4        | 4045.5        |
| 6           | 751.9        | 1471.7        | 3936.2        |
| 7           | 740.0        | 1446.1        | 3960.1        |
| <b>min</b>  | <b>735.4</b> | <b>1429.2</b> | <b>3936.2</b> |
| <b>mean</b> | <b>747.5</b> | <b>1456.5</b> | <b>4001.0</b> |
| <b>max</b>  | <b>759.2</b> | <b>1476.6</b> | <b>4059.5</b> |

### node5502 (8 x B200)

| GPU | FP32  | BF16   | FP8    |
|-----|-------|--------|--------|
| 0   | 759.7 | 1445.1 | 4087.0 |
| 1   | 769.3 | 1441.1 | 4113.7 |
| 2   | 755.0 | 1434.7 | 4048.0 |
| 3   | 764.5 | 1447.6 | 4092.0 |
| 4   | 766.8 | 1444.8 | 4063.4 |

|             |              |               |               |
|-------------|--------------|---------------|---------------|
| 5           | 757.4        | 1432.0        | 4011.1        |
| 6           | 755.1        | 1424.9        | 4058.6        |
| 7           | 754.5        | 1425.6        | 4026.0        |
| <b>min</b>  | <b>754.5</b> | <b>1424.9</b> | <b>4011.1</b> |
| <b>mean</b> | <b>760.3</b> | <b>1437.0</b> | <b>4062.5</b> |
| <b>max</b>  | <b>769.3</b> | <b>1447.6</b> | <b>4113.7</b> |

Converged = the final sustained-average throughput gpu-fryer reports per GPU at the end of each precision run. Higher is better; large spread across GPUs or any throttling flag indicates a problem.