

gpu-fryer summary

- Generated: 2026-07-13 16:24:13
- Nodes: node5500, node5502 (8 x NVIDIA B200 each)
- Precisions: FP32, BF16, FP8
- Reference (MIT aicr-benchmarks, gpu-fryer/summary.md, b0025): per-GPU mean TFLOP/s — FP32 772, BF16 1500, FP8 4115

Per-node mean converged throughput (TFLOP/s)

Node	FP32	BF16	FP8	Health
node5500	746	1454	3989	ok
node5502	751	1461	4011	ok
reference (b0025)	772	1500	4115	—

% of B200 reference (mean)

Node	FP32	BF16	FP8
node5500	97%	97%	97%
node5502	97%	97%	97%

Per-GPU converged throughput (TFLOP/s)

node5500

GPU	FP32	BF16	FP8
0	742.1	1476.0	4044.7
1	756.6	1470.8	4059.5
2	740.1	1437.4	3928.2
3	740.5	1437.4	3940.0
4	757.1	1466.0	4006.7
5	749.9	1456.6	4030.8
6	745.1	1446.8	3920.7
7	740.6	1442.1	3978.9

min	740.1	1437.4	3920.7
mean	746.5	1454.1	3988.7
max	757.1	1476.0	4059.5

node5502

GPU	FP32	BF16	FP8
0	752.1	1468.6	4033.6
1	759.3	1467.5	4069.7
2	746.2	1460.9	3987.1
3	758.0	1470.4	4040.0
4	755.9	1471.1	4022.1
5	746.7	1453.7	3955.5
6	744.9	1448.1	4003.8
7	744.4	1449.5	3980.1
min	744.4	1448.1	3955.5
mean	751.0	1461.2	4011.5
max	759.3	1471.1	4069.7

Converged = the final sustained-average throughput gpu-fryer reports per GPU at the end of each precision run. Higher is better; large spread across GPUs or any throttling flag indicates a problem.